

Best run per LLM vs full-dataset baselines (T5-Large, Plan-Simp)
BLEU

